

Support Vector Machines (SVMs)

References:


Cristianini & Shawe-Taylor book; Vapnik's book; and
“A Tutorial on Support Vector Machines for Pattern
Recognition” by Burges, in *Datamining and Knowledge
Discovery 2, 1998 (check citeseer)*
kernel-machines.org

SVMs

- Combines learning and optimization theory
- *Exactly* solve optimization (as opposed to ANN or Decision Trees)
- Allows “kernel trick” to get more features almost for free
- Soft margin relaxes linear separable assumption


Basic Idea: (separable case)

- Which linear threshold hypothesis to use?
- Pick one with biggest *margin* (gap)
 - Intuitively good, safe against small errors
 - Some supporting theory
 - Empirically works well
- Leads to a quadratic programming problem
(Vapnik 95, Cortes and Vapnik 95)



Basic Idea:

- Pick one with biggest *margin* (gap) δ
- Leads to a quadratic programming problem
(Vapnik 95, Cortes and Vapnik 95)
- **Support Vectors** are the points with smallest margin (closest to boundary, circled)
- hypothesis stability: it only changes if support vectors change



Good generalization

- If $f(\mathbf{x}) = \sum w_j x_j + b$ and $\sum w_j^2 = 1$ and large margin, then it (expects) to generalize well

generalization error: $\tilde{O} \left(\frac{R^2}{N\delta^2} + \frac{\log(1/\text{conf})}{N} \right)$

$R = \text{radius of support}$, $N = \# \text{ examples}$, $\delta = \text{margin}$

*this is independent of dimension
(# of features) !*

SVM applications

- Text classification - (Joachims, used outlier removal when a_i too large, Joachims also author of SVM-light)
- Object detection (e.g. Osuna et.al., for faces)
- Hand written digits - few support vectors
- Bioinformatics - protein homology, micro-array
- May be best current learning method

See Cristianini&Shaw-Taylor book, and

[www. kernel-machines.org](http://www.kernel-machines.org) for more info and references

SVM Mathematically

- Find \mathbf{w} , b , δ such that:

$$\mathbf{w} \cdot \mathbf{x}_i + b \geq \delta \quad \text{when } y_i = +1$$

$$\mathbf{w} \cdot \mathbf{x}_i + b \leq -\delta \quad \text{when } y_i = -1$$

and δ as big as possible

- Scaling issue: fix by setting $\delta=1$ and finding shortest \mathbf{w} :

$$\min_{\mathbf{w}, b} \|\mathbf{w}\|^2 \text{ subject to}$$

$$y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1 \quad \text{for all examples } (\mathbf{x}_i, y_i)$$


- This is a quadratic programming problem, packages exist (also SMO algorithm)

More Math **not so** (Quickly)

- LaGrange multipliers, dual problem, ..., magic, ...
- The \mathbf{w} vector is a linear combination of support vectors: $\mathbf{w} = \sum_{\text{sup vects } i} a_i \mathbf{x}_i$
- Predict on \mathbf{x} based on $\mathbf{w} \cdot \mathbf{x} \geq b$, or equivalently $\sum_{\text{sup vects } i} a_i (\mathbf{x}_i \cdot \mathbf{x}) \geq b$
- **Key idea 1:** predictions (and finding a_i 's) depend only on dot products (also true for least squares and perceptron)


Soft Margin for noise

- Allow margin errors
- $\varepsilon_i \geq 0$ measures amount of error on \mathbf{x}_i
- Want to minimize both $\mathbf{w} \cdot \mathbf{w}$ and $\sum_i \varepsilon_i$
- Need tradeoff parameter
- Still Quadratic programming, math messier





Kernel functions

- Don't need to use dot product, can use any “dot-product like” function $K(\mathbf{x}, \mathbf{x}')$
- Example: $K(\mathbf{x}, \mathbf{x}') = (\mathbf{x} \cdot \mathbf{x}' + 1)^2$
- In 1-dimension: $K(x, x') = x^2 x'^2 + 2xx' + 1$
 $= (x^2\sqrt{2} \quad x, 1) \cdot (x'^2\sqrt{2} \quad x', 1)$, get quadratic term too (extra feature)
- How can this help?



Want to classify these points with a SVM,
No hyper-plane (threshold) has good margin



Kernel trick:

- Kernel functions embed the low dimensional original space into a higher dimensional one, like the film of a soap bubble
- Although sample is not linearly separable in the original space, the higher-dimensional embedding might be
- Get embedding and “extra” features for free (almost)
- General trick for any dot-product based algorithm (e.g. Perceptron)

Common Kernels

- Polynomials: $K(\mathbf{x}, \mathbf{x}') = (\mathbf{x} \cdot \mathbf{x}' + 1)^d$
 - Gets new features like x_1x_3
- Radial Basis function (Gaussian):
 $K(\mathbf{x}, \mathbf{x}') = \exp(-\|\mathbf{x} - \mathbf{x}'\|^2 / 2\sigma^2)$
- Sigmoid: $K(\mathbf{x}, \mathbf{x}') = \tanh(a (\mathbf{x} \cdot \mathbf{x}') - b)$
- Also special purpose string and graph kernels

Which to use and parameters? cross validation can help, RBF Kernels powerful

What can be a Kernel function?

- $K(\mathbf{x}, \mathbf{z})$ is a (Mercer) kernel function if
 $K(\mathbf{x}, \mathbf{z}) = \phi(\mathbf{x}) \cdot \phi(\mathbf{z})$ implies symmetric
for some feature transformation $\phi()$
- $K(\mathbf{x}, \mathbf{z})$ is a Mercer kernel function if the matrix $\mathbf{K} = (k(\mathbf{x}_i, \mathbf{x}_j))_{i,j}$ is positive semi-definite for all subsets $\{\mathbf{x}_1, \dots, \mathbf{x}_n\}$ of instance space
- Vapnik showed expected test error bound:
(expected # support vectors) / (# examples)

Gaussian (RBF) Kernels

- Act like weighted nearest neighbor on support vectors
- Allow very flexible hypothesis space
- Mathematically a dot product in an “infinite dimensional” space
- SVMs unify LTU (linear kernel) and nearest neighbor (Gaussian kernel)

recall $\mathbf{w} = \sum \alpha_i \mathbf{x}_i$ what is the prediction?

Multiclass Kernel Machines

- 1-vs-all
- Pairwise separation (all pairs)
- Error-Correcting Output Codes
- Single multiclass optimization

$$\min \frac{1}{2} \sum_{i=1}^K \|\mathbf{w}_i\|^2 + C \sum_i \sum_t \xi_i^t$$

subject to

t superscript = example #
z^{*t*} = label of example *t*
i = class #

$$\mathbf{w}_{z^t}^T \mathbf{x}^t + w_{z^t 0} \geq \mathbf{w}_i^T \mathbf{x}^t + w_{i0} + 2 - \xi_i^t, \forall i \neq z^t, \xi_i^t \geq 0$$

SVM for Regression

Use a linear model (possibly kernelized)

$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

Use the ε -(in)sensitive error function


$$E_{\varepsilon}(y, f(\mathbf{x})) = \begin{cases} 0 & \text{if } |y - f(\mathbf{x})| < \varepsilon \\ |y - f(\mathbf{x})| - \varepsilon & \text{otherwise} \end{cases}$$


$$\min \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_t (\xi_+^t + \xi_-^t)$$

$$y_n - (\mathbf{w}^T \mathbf{x}_n + w_0) \leq \varepsilon + \xi_n^+$$

$$(\mathbf{w}^T \mathbf{x}_n + w_0) - y_n \leq \varepsilon + \xi_n^-$$


$$\xi_n^-, \xi_n^+ \geq 0$$



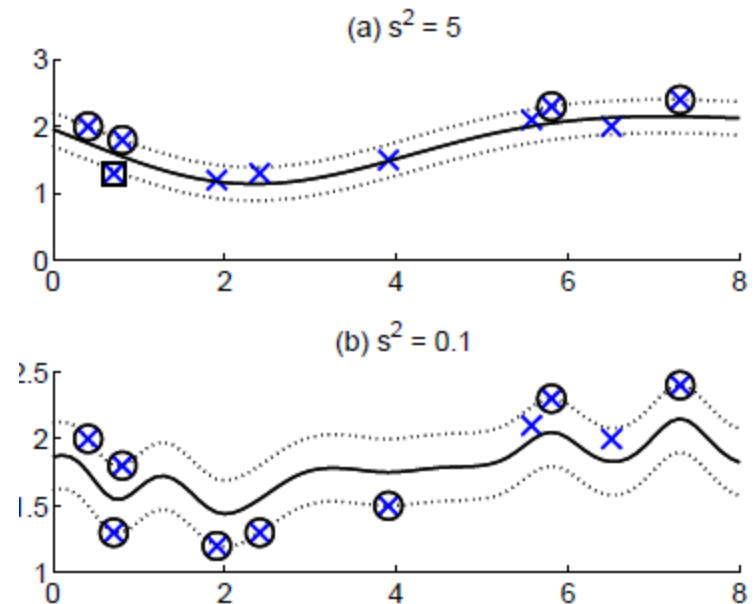


Kernel Regression

- Polynomial kernel



- Gaussian kernel



One-Class Kernel Machines

- Consider a sphere with center \mathbf{a} and radius R

$$\min R^2 + C \sum_t \xi^t$$


subject to

$$\|\mathbf{x}^t - \mathbf{a}\| \leq R^2 + \xi^t, \quad \xi^t \geq 0$$


$$L_d = \sum_t \alpha^t \left(x^t \right)^T x^t - \sum_{t=1}^N \sum_s \alpha^t \alpha^s r^t r^s \left(x^t \right)^T x^s$$

subject to

$$0 \leq \alpha^t \leq C, \quad \sum_t \alpha^t = 1$$



One-class SVM with Gaussian kernels (s is sigma)



SVM Summary

- Model: very flexible, but must pick many parameters (kernel, kernel parameters, trade-off)
- Data: Numeric (depending on kernel)
- Interpretable? Yes for dot product, pretty pictures for Gaussian kernels in low dimensions
- Missing values? No
- Noise/outliers? Very good (softmargin)
- Irrelevant features? Yes for dot product with abs value penalty
- Comp. efficiency? quadratic in # examples, but “exact” optimization rather than approximate (like ANN, decision trees)
Chunking techniques and other optimizations (SMO iteratively optimize a_i' s over pairs, SGD, PEGASOS, etc.)